

# Guillaume Payeur (260929164)

Note to grader: all the code explanations and figures are included in this notebook

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
import warnings
from scipy import interpolate
%matplotlib inline
import matplotlib as mpl
mpl.rcParams['figure.dpi'] = 200
plt.rcParams.update({"text.usetex": True})
```

## Q1

We find expressions for the error on the numerical derivatives computed using  $f(x \pm \delta)$  and using  $f(x \pm 2\delta)$ . For  $f(x \pm \delta)$  we have

$$\begin{aligned} f'(x) &= \frac{f(x + \delta) - f(x - \delta)}{2\delta} \\ &= \frac{(f(x) + f'(x)\delta + \frac{1}{2}f''(x)\delta^2 + \dots)(1 + \epsilon_{g0}) - (f(x) - f'(x)\delta + \frac{1}{2}f''(x)\delta^2 - \dots)(1 + \epsilon_{g1})}{2\delta} \\ &= f'(x) + \frac{f(x)\epsilon + \frac{1}{3}f'''(x)\delta^3 + \frac{1}{60}f^{(5)}(x)\delta^5}{2\delta} \end{aligned}$$

so the error on the derivative is (ignoring negligible terms)

$$E = \frac{1}{2} \frac{f\epsilon}{\delta} + \frac{1}{6} f''' \delta^2 + \frac{1}{120} f^{(5)} \delta^4 \quad (4)$$

Doing the same for  $f(x \pm 2\delta)$  we get that the error is

$$E = \frac{1}{4} \frac{f\epsilon}{\delta} + \frac{2}{3} f''' \delta^2 + \frac{2}{15} f^{(5)} \delta^4 \quad (5)$$

What we want is to take a combination of the two computed derivatives that cancels off the  $f'''$  term in the errors. That combination is

$$f'(x) = \frac{1}{3} \left( 4 \left( \frac{f(x + \delta) - f(x - \delta)}{2\delta} \right) - \left( \frac{f(x + 2\delta) - f(x - 2\delta)}{4\delta} \right) \right) \quad (6)$$

in which case the error is

$$E \sim \frac{f\epsilon}{\delta} + f^{(5)} \delta^4 \quad (7)$$

Minimizing with respect to  $\delta$  we get that the optimal  $\delta$  is

$$\delta = \left( \frac{f\epsilon}{f^{(5)}} \right)^{1/4} \quad (8)$$

In [2]:

```
# Numerical derivative
def num_derivative(a,x,delta):
    # derivative using f(x +/- delta)
    f_prime_1 = (np.exp(a*(x+delta))-np.exp(a*(x-delta)))/(2*delta)
    # derivative using f(x +/- 2delta)
    f_prime_2 = (np.exp(a*(x+2*delta))-np.exp(a*(x-2*delta)))/(4*delta)
    return (4*f_prime_1-f_prime_2)/3

# Analytic (true) derivative
def true_derivative(a,x):
    return a*np.exp(a*x)
```

Testing for  $f(x) = e^x$  at  $x = 0$ . Here we can assume  $f \sim f^{(5)}$  and the boxed equation for optimal  $\delta$  gives


$$\delta = \left( 10^{-16} \right)^{1/4} = 10^{-4} \quad (9)$$

In [3]:

```
x=0
a=1
delta_array = 10**(np.linspace(-8,0,100))
error_array = np.zeros_like(delta_array)
for i in range(delta_array.shape[0]):
    error_array[i] = np.abs(num_derivative(a,x,delta_array[i])-true_derivative(a,x))
```

In [4]:

```
plt.plot(delta_array,error_array,color='green')
plt.plot([1e-4,1e-4],[error_array.min(),error_array.max()],ls='--',color='gray',label='
plt.xscale('log')
plt.yscale('log')
plt.xlabel('$\delta$',fontsize=15)
plt.ylabel('Absolute error on derivative',fontsize=15)
plt.title('Error on derivative for $f(x)=e^x$',fontsize=20)
plt.legend(loc=0,frameon=True,shadow=True,fontsize=15)
plt.tight_layout()
```




Testing for  $f(x) = e^{0.01x}$  at  $x = 0$ . Here we can assume  $f \sim 10^2 f^{(5)}$  and the boxed equation for optimal  $\delta$  gives

$$\delta = \left(10^{-16}(10^2)^5\right)^{1/4} = 10^{-3/2} \quad (10)$$

```
In [5]:
x=1
a=0.01
delta_array = 10**(np.linspace(-8,0,100))
error_array = np.zeros_like(delta_array)
for i in range(delta_array.shape[0]):
    error_array[i] = np.abs(num_derivative(a,x,delta_array[i])-true_derivative(a,x))
```

```
In [6]:
plt.plot(delta_array,error_array,color='green')
plt.plot([0.03,0.03],[error_array.min(),error_array.max()],ls='--',color='gray',label='
plt.xscale('log')
plt.yscale('log')
plt.xlabel('$\delta$',fontsize=15)
plt.ylabel('Absolute error on derivative',fontsize=15)
plt.title('Error on derivative for $f(x)=e^{\{0.01x\}}$',fontsize=20)
plt.legend(loc=0,frameon=True,shadow=True,fontsize=15)
plt.tight_layout()
```



The agreement between predictions and observations isn't stellar but that's because I've made big approximations when calculating the interpolation error: I assumed that all coefficients in the expression for  $E$  are 1.

## Q2

In order to find the optimal  $\delta$  to minimize the error on the derivative, we need an estimate for  $(f/f''')$ , so we need to compute an estimate of the third derivative. We can do this by calculating

$$f'''(x) = \frac{1}{2\delta}(f''(x+\delta) - f''(x-\delta)) \quad (11)$$

now plugging in

$$f''(x) = \frac{1}{2\delta}(f'(x+\delta) - f'(x-\delta)) \quad (12)$$

and then plugging in

$$f'(x) = \frac{1}{2\delta}(f(x+\delta) - f(x-\delta)) \quad (13)$$

we get

$$f'''(x) = \frac{1}{2\delta}(f(x+3\delta) - 3f(x+\delta) + 3f(x-\delta) - f(x-3\delta)) \quad (14)$$

So we use this equation with  $\delta = 10^{-5}$  (we don't care about  $\delta$  so much here since we just care about the order of magnitude of  $f'''$ ) and then use the equation derived in class, namely

$$\delta = \left( \frac{\epsilon f}{f'''} \right)^{1/3} \quad (15)$$

to find the optimal  $\delta$ . Then we compute the derivative via

$$f'(x) = \frac{1}{2\delta} (f(x + \delta) - f(x - \delta)) \quad (16)$$

We return this value of  $f'(x)$  aswell as  $\delta$  if requested. As for the size of the error, we learned in class that the error here is roughly

$$E = \frac{\epsilon f}{\delta} + 2f'''\delta^2 \quad (17)$$

and that's what we return as error estimate if requested.

```
In [7]: def ndiff(fun,x,full=False):
    eps = 1e-16
    delta_f3 = 1e-5

    # Approximating third derivative
    f = fun(x)
    f3 = 1/(8*delta_f3**3)*(fun(x+3*delta_f3)-3*fun(x+delta_f3)+3*fun(x-delta_f3)-fun(x-3*delta_f3))
    delta = (f*eps/f3)**(1/3)

    # Calculating derivative
    f1 = 1/(2*delta)*(fun(x+delta)-fun(x-delta))
    err = eps*f/delta + 2*f3*delta**2

    if full:
        return f1,delta,err
    return f1
```

Now we just do a small test of ndiff. In each case, I take a derivative, and compare both the numerical derivative and the estimated error to the truth.

Testing with  $f(x) = e^x$  and  $x = 0$

```
In [8]: fun=np.exp
fprime,delta,e_fprime = ndiff(fun,0,full=True)
print('dx:',delta)
print('derivative:',fprime)
print('true derivative:', 1)
print('estimated error on derivative:',e_fprime)
print('true error on derivative:',np.abs(fprime-1))
```

```
dx: 4.621528465628242e-06
derivative: 1.0000000000030995
true derivative: 1
estimated error on derivative: 6.491358913640684e-11
true error on derivative: 3.099520640148512e-12
```

Testing with  $f(x) = e^x$  and  $x = 1$

In [9]:

```

fun=np.exp
fprime,delta,e_fprime = ndiff(fun,1,full=True)
print('dx:',delta)
print('derivative:',fprime)
print('true derivative:', np.e)
print('estimated error on derivative:',e_fprime)
print('true error on derivative:',np.abs(fprime-np.e))

```

```

dx: 4.739350250110249e-06
derivative: 2.71828182848983
true derivative: 2.718281828459045
estimated error on derivative: 1.7206674027072477e-10
true error on derivative: 3.078470811601619e-11

```

Testing with  $f(x) = e^{0.1x}$  and  $x = 0$

In [10]:

```

def fun(x):
    return np.e**(0.1*x)
fprime,delta,e_fprime = ndiff(fun,0,full=True)
print('dx:',delta)
print('derivative:',fprime)
print('true derivative:',0.1)
print('estimated error on derivative:',e_fprime)
print('true error on derivative:',np.abs(fprime-0.1))

```

```

dx: 1.2167646631985893e-05
derivative: 0.10000000000403747
true derivative: 0.1
estimated error on derivative: 2.4655548363096828e-11
true error on derivative: 4.03746480692746e-12

```

Testing with  $f(x) = e^{0.1x}$  and  $x = 1$

In [11]:

```

def fun(x):
    return np.e**(0.1*x)
fprime,delta,e_fprime = ndiff(fun,1,full=True)
print('dx:',delta)
print('derivative:',fprime)
print('true derivative:',0.1*np.e**0.1)
print('estimated error on derivative:',e_fprime)
print('true error on derivative:',np.abs(fprime-0.1*np.e**0.1))

```

```

dx: 1.5849895500790092e-05
derivative: 0.11051709180254779
true derivative: 0.11051709180756478
estimated error on derivative: 2.0918199454764066e-11
true error on derivative: 5.01698682597862e-12

```

The estimates of the derivative are good and so are the estimates of the error, so this worked!

## Q3

Since we are given  $dV/dT$ , and since the intervals between the data points are quite small, I do a linear interpolation from the data point closest to the value of  $V$  requested, that is, I compute

$$T(V) = T(V_0) + \frac{dT}{dV}(V - V_0) \quad (18)$$

As for the error, I observe that the values of  $dV/dT$  are accurate to roughly a part in 100 (based on significant figures), so I approximate the error on the slope as being


$$\delta(dT/dV) = \frac{1}{100} \frac{dT}{dV} \quad (19)$$

and get my approximate error on  $T(V)$  by propagating the error on the slope to  $T(V)$ . I get

$$\delta(T(V)) = \frac{1}{100} \frac{dT}{dV}(V - V_0) \quad (20)$$

Let's begin by plotting the raw data

```
In [12]: # Plotting the raw data
dat=np.loadtxt('lakeshore.txt')
V = dat[:,1]
T = dat[:,0]
plt.plot(V,T,lw=0,marker='.',color='black')
plt.xlabel('$V$',fontsize=15)
plt.ylabel('$T$',fontsize=15)
plt.tight_layout()
```



```
In [13]: def lakeshore(V,data):
# Extracting the Voltage, Temperature, derivative of Temperature wrt Voltage
data_V = data[:,1]
data_T = data[:,0]
data_dTdV = 1000/data[:,2]
```

```

# Making arrays that will contain the interpolated temperatures and errors
T_interp = np.zeros_like(V)
e_T_interp = np.zeros_like(V)

for i in range(V.shape[0]):
    # Extracting the Voltage, Temperature, derivative of Temperature wrt Voltage
    data_V = data[:,1]
    data_T = data[:,0]
    data_dTdV = 1000/data[:,2]

    # If the requested value of V requires extrapolation, give a warning and return
    V_max = data_V[0]
    V_min = data_V[-1]
    if V[i] > V_max or V[i] < V_min:
        warnings.warn('A value of V requested is not within the range of data recorded')
        return np.nan, np.nan

    # Finding the index of Voltage data point closest to V
    V_right = np.where(data_V < V[i])[0][0]
    V_left = np.where(data_V > V[i])[0][-1]
    if np.abs(V_left - V[i]) < np.abs(V_right - V[i]):
        V_index = V_left
    else:
        V_index = V_right

    # Now doing the linear interpolation
    T_interp[i] = data_T[V_index] + data_dTdV[V_index]*(V[i]-data_V[V_index])

    # Getting the error estimate
    e_T_interp[i] = (data_dTdV[V_index]/100)*(V[i]-data_V[V_index])

# Returning the interpolated value and error
return T_interp, e_T_interp

```

Now I do a quick test where I evaluate  $T(V=0.5)$ .

```

In [14]: data=np.loadtxt('lakeshore.txt')
          lakeshore(np.array([0.5]),data)

```

```

Out[14]: (array([325.75905172]), array([0.05759052]))

```

I plot  $T(V=0.5)$  along with neighboring points to check that it worked


```

In [15]: data=np.loadtxt('lakeshore.txt')
          V=np.array([0.5])
          T_interp,e_T_interp = lakeshore(V,data)
          T_interp = T_interp[0]
          e_T_interp[e_T_interp>0]

          plt.plot(data[118:135,1],data[118:135,0],lw=0,marker='.',label='data',color='black')
          plt.plot(V,T_interp,marker='.',color='green',ms=10,lw=0,label='interpolation at $V=0.5$')
          plt.legend(loc=0,frameon=True,shadow=True,fontsize=15)
          plt.xlabel('$V$',fontsize=15)
          plt.ylabel('$T$',fontsize=15)
          plt.tight_layout()

```





And just as proof that my function can take in an array of  $V$ :

```
In [16]: data=np.loadtxt('lakeshore.txt')
lakeshore(np.array([0.5,0.8]),data)
```

```
Out[16]: (array([325.75905172, 192.44953271]), array([0.05759052, 0.02449533]))
```

## Q4

I adapt code from Professor Sievers to do polynomial interpolations, cubic spline interpolations, and rational fits. I take the function  $f(x) = \cos(x)$  with 13 points uniformly sampled from the interval  $(-\pi/2, \pi/2)$ .

Then I make a plot of the absolute interpolation error from all three approaches

```
In [17]: xi=np.linspace(-np.pi/2,np.pi/2,13)
yi=np.cos(xi)

x=np.linspace(xi[1],xi[-2],1001)
y_true=np.cos(x)

# Polynomial (cubic) interpolation
def poly(xi,yi,x):
    y_interp_poly=np.zeros(len(x))
    for i in range(len(x)):
        ind=np.max(np.where(x[i]>=xi)[0])
        x_use=xi[ind-1:ind+3]
        y_use=yi[ind-1:ind+3]
```

```

    pars=np.polyfit(x_use,y_use,3)
    pred=np.polyval(pars,x[i])
    y_interp_poly[i]=pred

    return y_interp_poly

# Cubic spline
def spline(xi,yi,x):
    spln=interpolate.splrep(xi,yi)
    y_interp_spline=interpolate.splev(x,spln)

    return y_interp_spline

# Rational fit
def rat(xi,yi,x):
    n=int((xi.shape[0]-1)/2)
    m=n

    pcols=[xi**k for k in range(n+1)]
    pmat=np.vstack(pcols)

    qcols=[-xi**k*yi for k in range(1,m+1)]
    qmat=np.vstack(qcols)
    mat=np.hstack([pmat.T,qmat.T])
    coeffs=np.linalg.inv(mat)@yi

    num=np.polyval(np.flipud(coeffs[:n+1]),x)
    denom=1+x*np.polyval(np.flipud(coeffs[n+1:]),x)
    y_interp_rat=num/denom

    return y_interp_rat

y_interp_poly, y_interp_spline, y_interp_rat = poly(xi,yi,x), spline(xi,yi,x), rat(xi,y

```

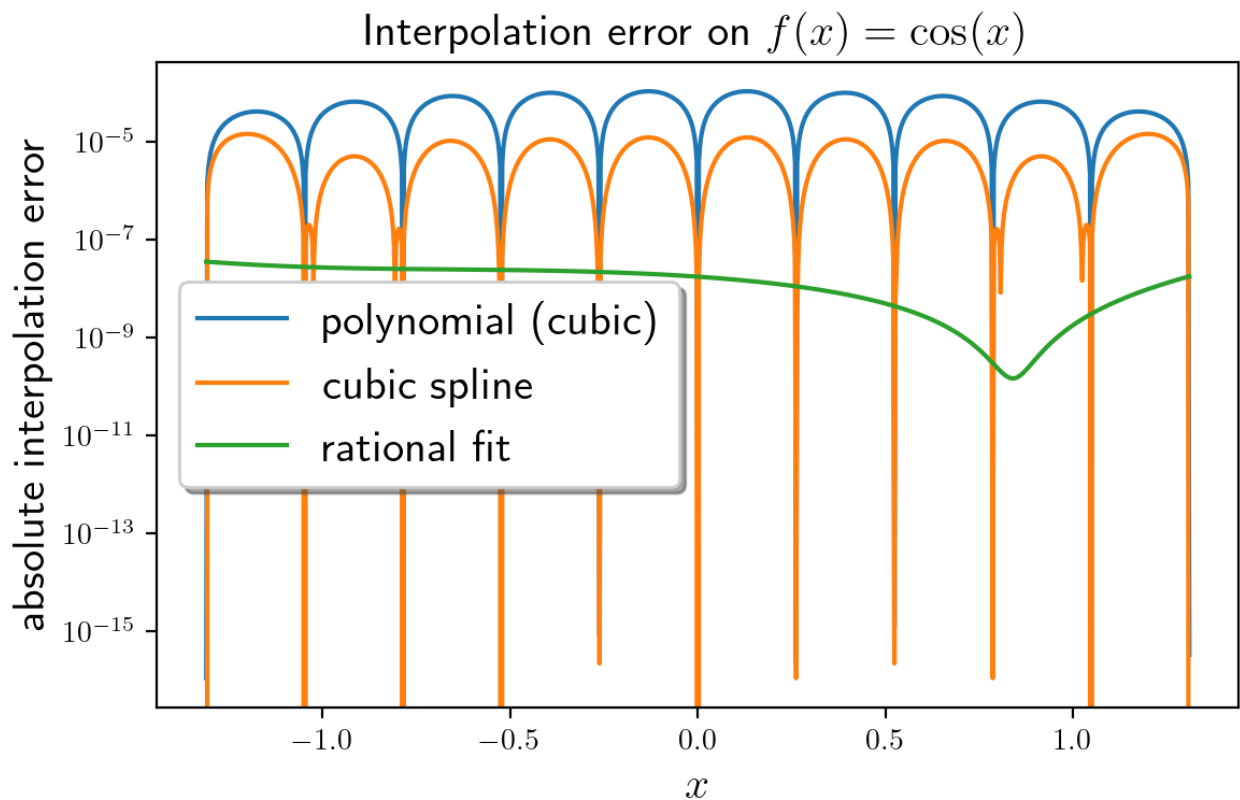
C:\Users\Guill\Anaconda3\envs\ml\_pytorch\lib\site-packages\ipykernel\_launcher.py:46: RankWarning: Polyfit may be poorly conditioned

In [18]:

```

plt.clf()
plt.plot(x,np.abs(y_interp_poly-y_true),label='polynomial (cubic)')
plt.plot(x,np.abs(y_interp_spline-y_true),label='cubic spline')
plt.plot(x,np.abs(y_interp_rat-y_true),label='rational fit')
plt.yscale('log')
plt.xlabel('$x$',fontsize=15)
plt.ylabel('absolute interpolation error',fontsize=15)
plt.title('Interpolation error on $f(x)=\cos(x)$',fontsize=15)
plt.legend(loc=0,frameon=True,shadow=True,fontsize=15)
plt.tight_layout()

```



What we see is that the rational fit was the most succesful by far, followed by the cubic spline.

Now I change the function to  $f(x) = \frac{1}{1+x^2}$  and do the same thing, this time with 5 points in the interval  $(-1, 1)$


```
In [19]: xi=np.linspace(-1,1,5)
yi=1/(1+xi**2)

x=np.linspace(xi[1],xi[-2],1001)
y_true=1/(1+x**2)

y_interp_poly, y_interp_spline, y_interp_rat = poly(xi,yi,x), spline(xi,yi,x), rat(xi,y
```

C:\Users\Guill\Anaconda3\envs\ml\_pytorch\lib\site-packages\ipykernel\_launcher.py:7: Rank Warning: Polyfit may be poorly conditioned  
import sys

```
In [20]: plt.clf()
plt.plot(x,np.abs(y_interp_poly-y_true),label='polynomial (cubic)')
plt.plot(x,np.abs(y_interp_spline-y_true),label='cubic spline')
plt.plot(x,np.abs(y_interp_rat-y_true),label='rational fit')
plt.yscale('log')
plt.xlabel('$x$',fontsize=15)
plt.ylabel('absolute interpolation error',fontsize=15)
plt.title('Interpolation error on Lorentzian with $5$ data points',fontsize=15)
plt.legend(loc=0,frameon=True,shadow=True,fontsize=15)
plt.tight_layout()
```



As expected the error on the rational fit is roughly machine precision ( $\sim 10^{-16}$ ) which was expected since the rational function can in principle fit the Lorentzian exactly (zero error). The cubic spline performs second best again.

Now I increase the number of points to 11 and see that the rational fit fails

```
In [21]: xi=np.linspace(-1,1,11)
yi=1/(1+xi**2)

x=np.linspace(xi[1],xi[-2],1001)
y_true=1/(1+x**2)

# Rational fit (modified to print p,q, and the determinant of the matrix we're invertin
def rat(xi,yi,x):
    n=int((xi.shape[0]-1)/2)
    m=n

    pcols=[xi**k for k in range(n+1)]
    pmat=np.vstack(pcols)

    qcols=[-xi**k*yi for k in range(1,m+1)]
    qmat=np.vstack(qcols)
    mat=np.hstack([pmat.T,qmat.T])
    coeffs=np.linalg.inv(mat)@yi

    num=np.polyval(np.flipud(coeffs[:n+1]),x)
    denom=1+x*np.polyval(np.flipud(coeffs[n+1:]),x)
    y_interp_rat=num/denom

    print('p:',coeffs[:n+1])
    print('q:',coeffs[n+1:])
    print('determinant:',np.linalg.det(mat))
```

```
return y_interp_rat
```

```
y_interp_poly, y_interp_spline, y_interp_rat = poly(xi,yi,x), spline(xi,yi,x), rat(xi,y
```

```
p: [ 3.4254269  16.          4.          -4.          -2.6467993  -2.18149787]
```


```
q: [ 4.  0. -1.  8. -8.]
```

```
determinant: 1.4333226137647608e-54
```

```
C:\Users\Guill\Anaconda3\envs\ml_pytorch\lib\site-packages\ipykernel_launcher.py:30: Ran  
kWarning: Polyfit may be poorly conditioned
```

In [22]:

```
plt.clf()
plt.plot(x,np.abs(y_interp_poly-y_true),label='polynomial (cubic)')
plt.plot(x,np.abs(y_interp_spline-y_true),label='cubic spline')
plt.plot(x,np.abs(y_interp_rat-y_true),label='rational fit')
plt.yscale('log')
plt.xlabel('$x$',fontsize=15)
plt.ylabel('absolute interpolation error',fontsize=15)
plt.title('Interpolation error on Lorentzian with 11 data points (np.linalg.inv)',fon
plt.legend(loc=0,frameon=True,shadow=True,fontsize=15)
plt.tight_layout()
```



This is because the matrix we are inverting is almost singular (its determinant is  $\sim 10^{-54}$  as printed above) and `np.linalg.inv` fails to take the inverse properly. Printing `p` and `q` as above we see that they have not been set to their correct values. We should have had that `p=(1,0,0,0,0,0)` and `q=(0,1,0,0,0)` because corresponds identically the lorentzian.

But `np.linalg.pinv` is more stable and we can find the inverse correctly. So here I change `np.linalg.inv` to `np.linalg.pinv` and see that the rational fit performs as expected again.

In [23]:

```
xi=np.linspace(-1,1,11)
```

```

yi=1/(1+xi**2)

x=np.linspace(xi[1],xi[-2],1001)
y_true=1/(1+x**2)

# Rational fit (modified to use np.linalg.pinv instead of np.linalg.inv)
def rat(xi,yi,x):
    n=int((xi.shape[0]-1)/2)
    m=n

    pcols=[xi**k for k in range(n+1)]
    pmat=np.vstack(pcols)

    qcols=[-xi**k*yi for k in range(1,m+1)]
    qmat=np.vstack(qcols)
    mat=np.hstack([pmat.T,qmat.T])
    coeffs=np.linalg.pinv(mat)@yi

    num=np.polyval(np.flipud(coeffs[:n+1]),x)
    denom=1+x*np.polyval(np.flipud(coeffs[n+1:]),x)
    y_interp_rat=num/denom

    return y_interp_rat

y_interp_poly, y_interp_spline, y_interp_rat = poly(xi,yi,x), spline(xi,yi,x), rat(xi,y

```

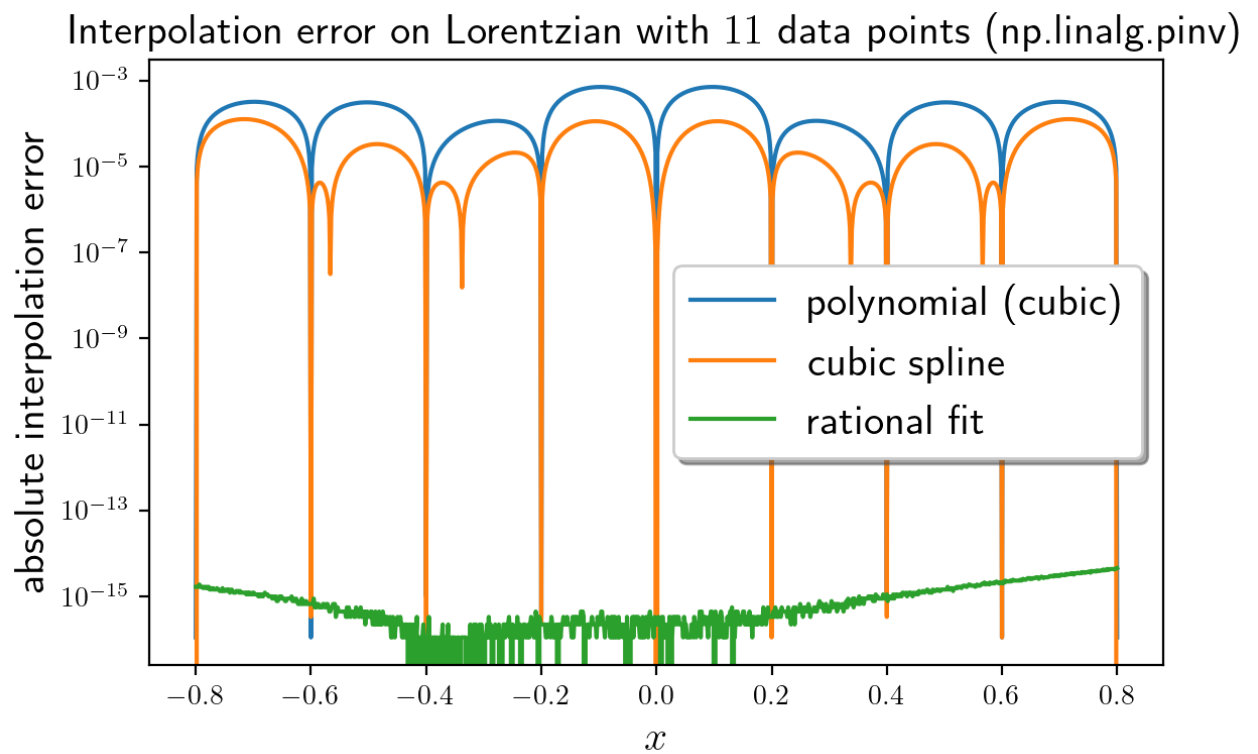
C:\Users\Guill\Anaconda3\envs\ml\_pytorch\lib\site-packages\ipykernel\_launcher.py:26: RankWarning: Polyfit may be poorly conditioned

In [24]:

```

plt.clf()
plt.plot(x,np.abs(y_interp_poly-y_true),label='polynomial (cubic)')
plt.plot(x,np.abs(y_interp_spline-y_true),label='cubic spline')
plt.plot(x,np.abs(y_interp_rat-y_true),label='rational fit')
plt.yscale('log')
plt.xlabel('$x$', fontsize=15)
plt.ylabel('absolute interpolation error', fontsize=15)
plt.title('Interpolation error on Lorentzian with $11$ data points (np.linalg.pinv)', fo
plt.legend(loc=0, frameon=True, shadow=True, fontsize=15)
plt.tight_layout()

```



The error on the rational fit is back to roughly machine precision, which is what was expected.